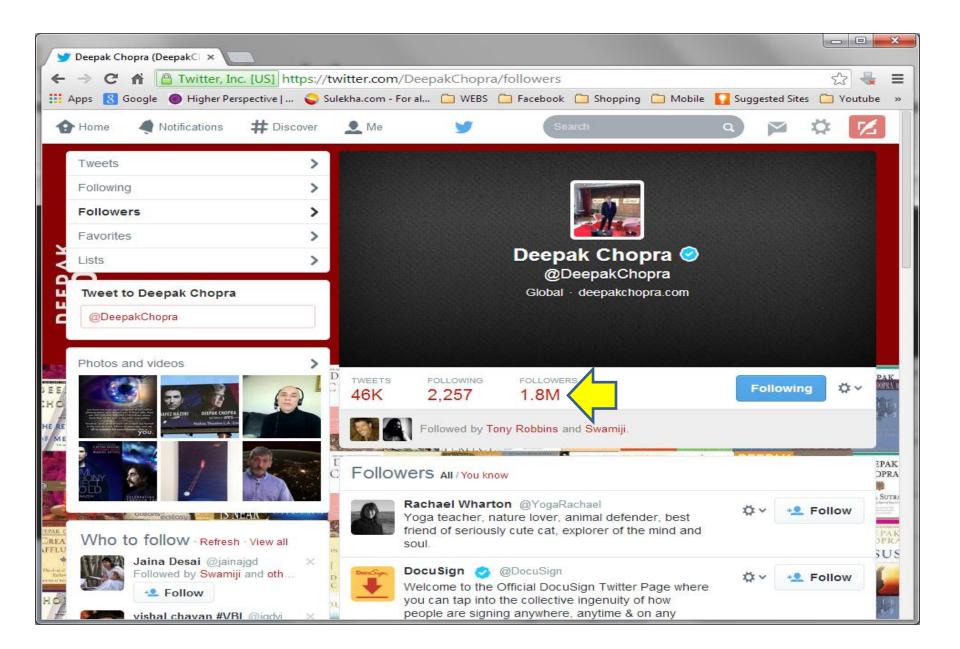
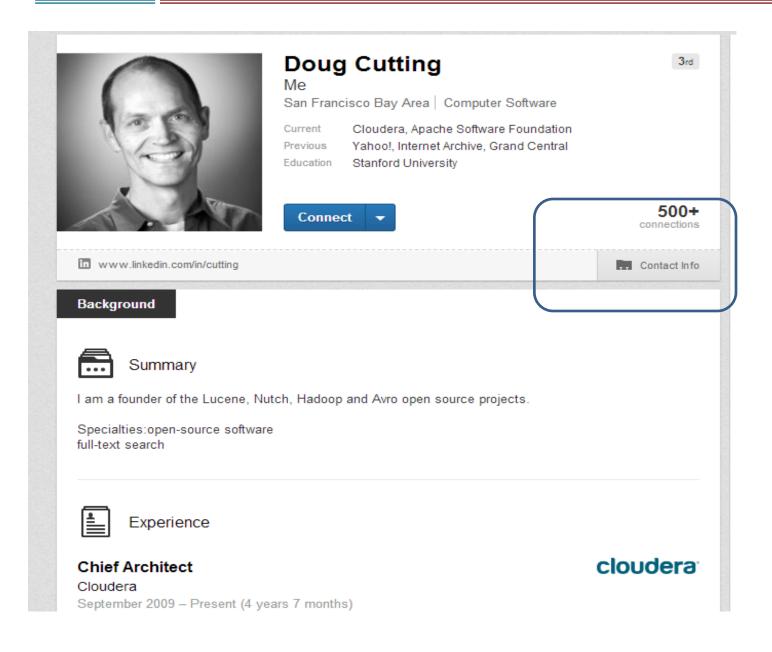
# **HBASE**

### Challenges in the Real World



## Challenges in the Real World



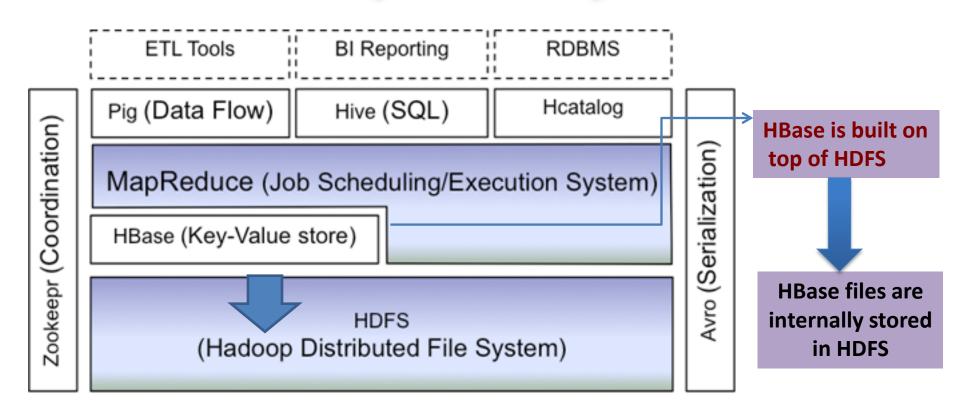
## Challenges in the Real World



- ✓ Fast Random Access
- ✓ Structured Data
- ✓ Huge Data
- ✓ Unstructured Data
- ✓ Variable Schema
- ✓ Need of Compression
- ✓ Need of the Distribution (Sharding)

- ✓ HBase is a distributed column-oriented data store built on top of HDFS
- ✓ HBase is an Apache open source project whose goal is to provide storage for the Hadoop Distributed Computing
- ✓ Data is logically organized into tables, rows and columns

# Hadoop Eco-System

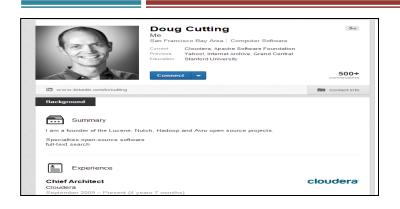


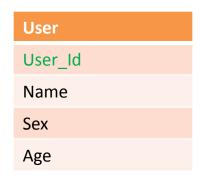
 Both HBase and HDFS are distributed systems that scale to hundreds or thousands of nodes

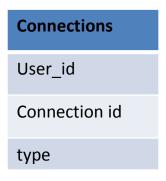
 HDFS is for batch processing (scans over big files) and does not support :

- for record lookup
- incremental addition of small batches
- Records level updates

# How tradition system will solve this?







User		
User ID	1	
Name	Doug Cutting	
Sex	M	
Age	45	

User Connections				
User-id	Connection-id	Туре		
1	1-2	Type1		
1	1-3	Type1		
1	1-4	Type1		
1	1-5	Type1		
1	1-6	Type1		
1	1-7	Type1		
	•			
1	1-500	Type1		

	Column Family 1 - User Details				Family 2 - ion-details
User ID	Name	Sex	Age	Connection-id- 01	Type-id-01
1	DC	M	45	1-2	Type1

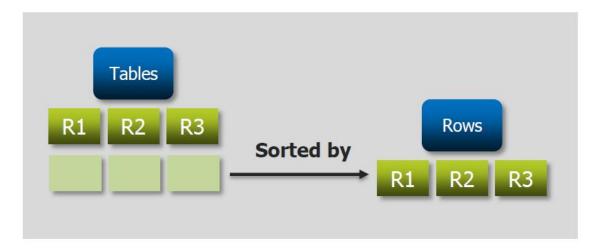
Column Family 1 - User Details			Column	Family : deta	2 - Connecti ails	ion-	
User ID	Name	Sex	Age	Connection-id- 01	Type-id- 01	Connection-id- 02	Type-id- 02
1	DC	M	45	1-2	Type1	1-3	Type1

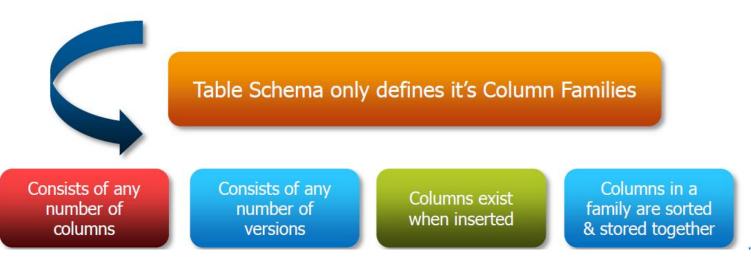


Colur 1 - Us											
User ID	Name	Sex	Age	Connection-id- 01	Type-id- 01	Connection-id- 02	Type-id- 02	Connection- id-03	Type-id- 03	Connection-id- 499	Type-id-499
1	DC	M	45	1-2	Type1	1-3	Type1	1-4	Type1	 1-499	Type1

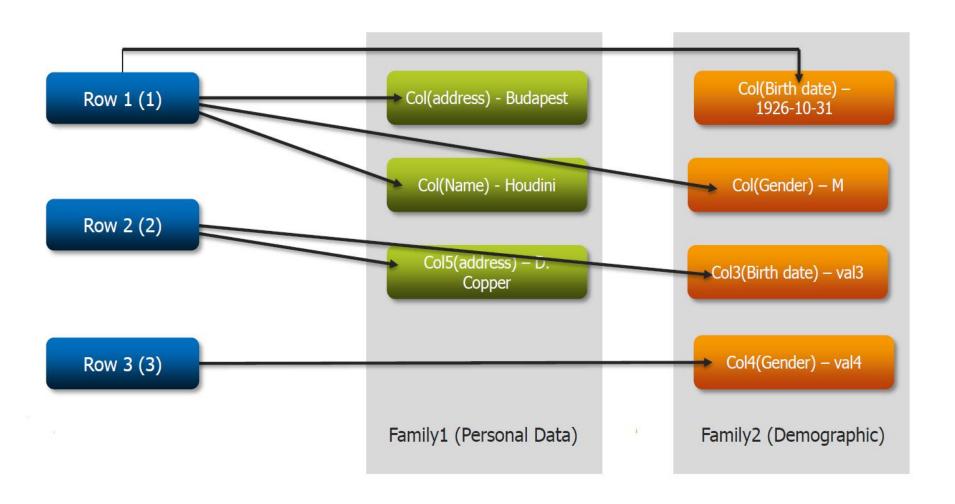
## Data Model

- Distributed Database
- Sorted Data
- Sparse Data Store
- Automatic sharding
- Multi dimensional
- Map
- ❖ Column = Column Family : Column Qualifier





Row key	Persona	l_data	demogra	phic
Persons ID	Name	Address	Birth Date	Gender
1	H. Houdini	Budapest	1926-10-31	М
2	D. Copper		1956-09-16	M
3	Merlin		1136-12-03	F
4				M
500,000,000	F. Cadillac	Nevada	1964-01-07	M



```
hbase(main):002:0> scan 'employee'

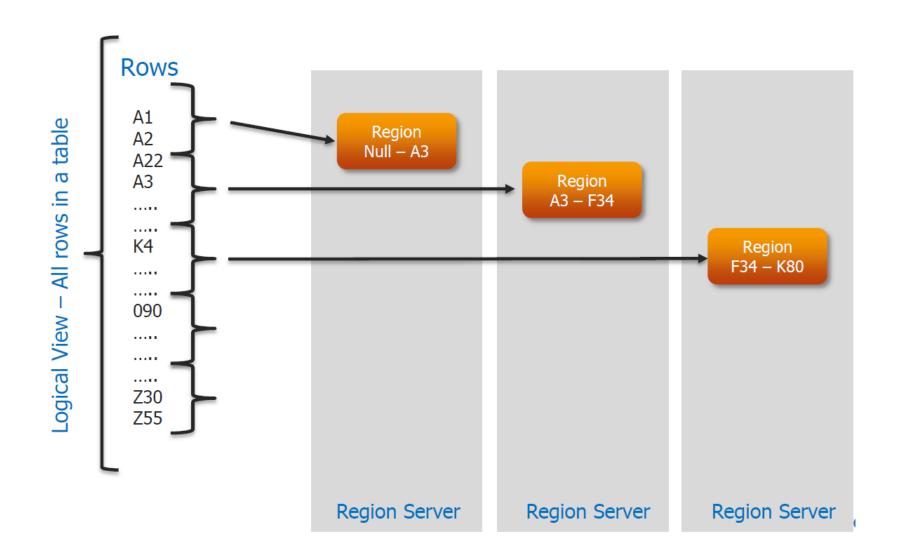
ROW COLUMN+CELL

emp1 column=address:city, timestamp=1396093072595, value=Mumbai

emp1 column=address:state, timestamp=1396093081699, value=MAH

emp1 column=address:street, timestamp=1396093010180, value=Thakur Village
```

Row Key	Column Family	Column qualifier	Values
Value- emp1	Value - address	Value – city, state, street	Value – Mumbai, MAH, Thakur Village
Unique for each row	<ul> <li>For Faster access</li> <li>less no of</li> <li>Column family</li> <li>are preferred.</li> </ul>	<ul> <li>Col qualifiers are NOT fixed. Can be added on the fly.</li> </ul>	<ul> <li>Various version of values are maintained</li> </ul>
<ul> <li>Identified each row</li> </ul>	CF are fixed		



Master



# Zookeeper

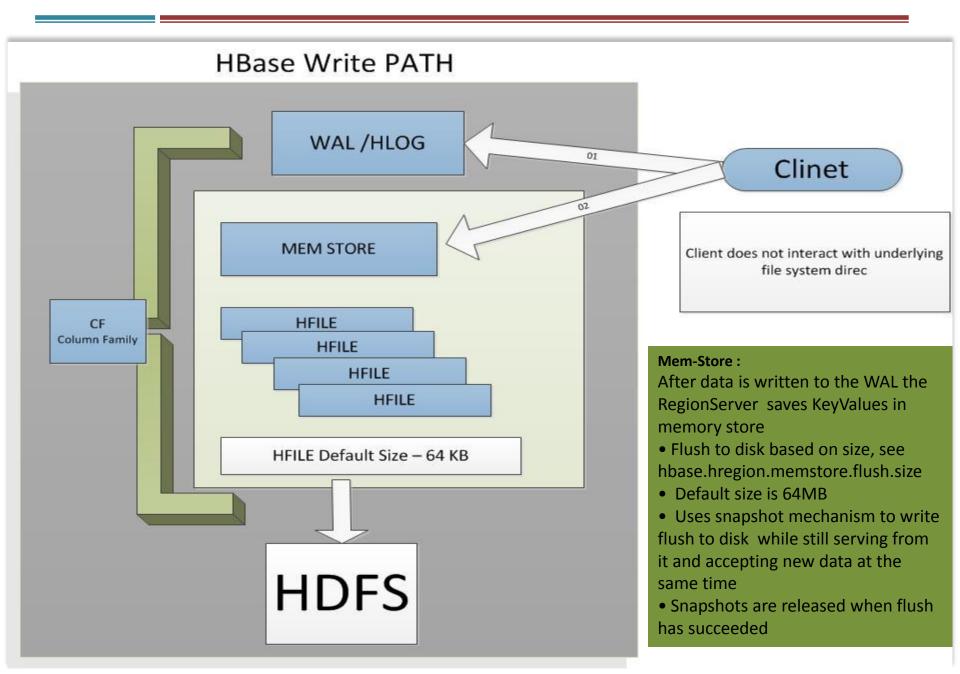
/hbase/region /hbase/region

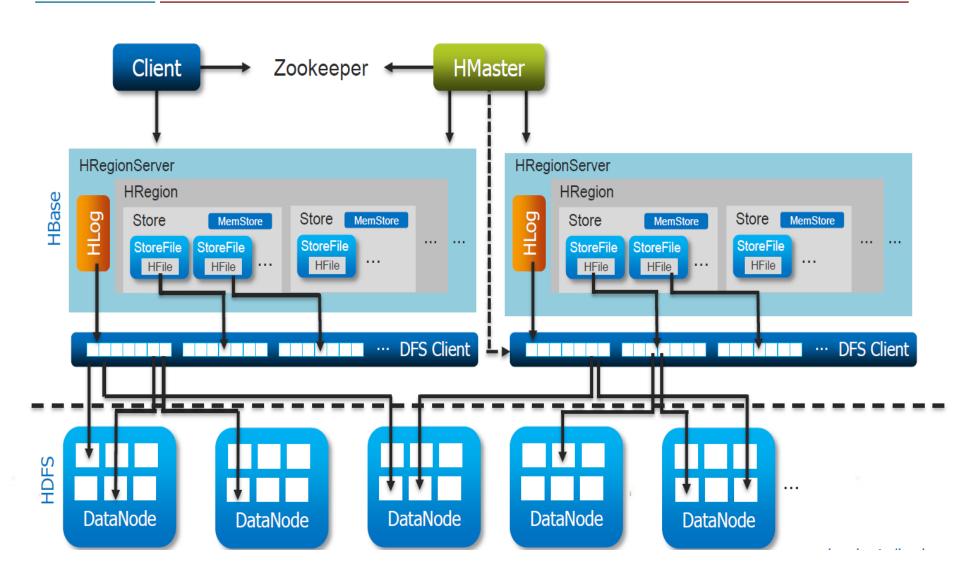
/hbase/region

**HDFS** 



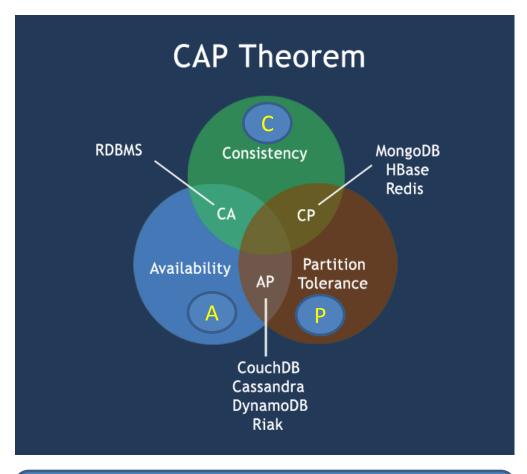
WAL





## Compactions

- ✓ HBase writes out immutable files as data is added
  - ✓ Each store consists of rowkey-ordered files
  - ✓ Immutable More files accumulate over time.
- ✓ Compaction rewrites several files into one.
  - ✓ Lesser files faster reads.
- ✓ Major compaction rewrites all files in a store into one.
  - ✓ Can drop deleted records and old versions.
- ✓ In a minor compaction, files to compact are selected based on a heuristic.

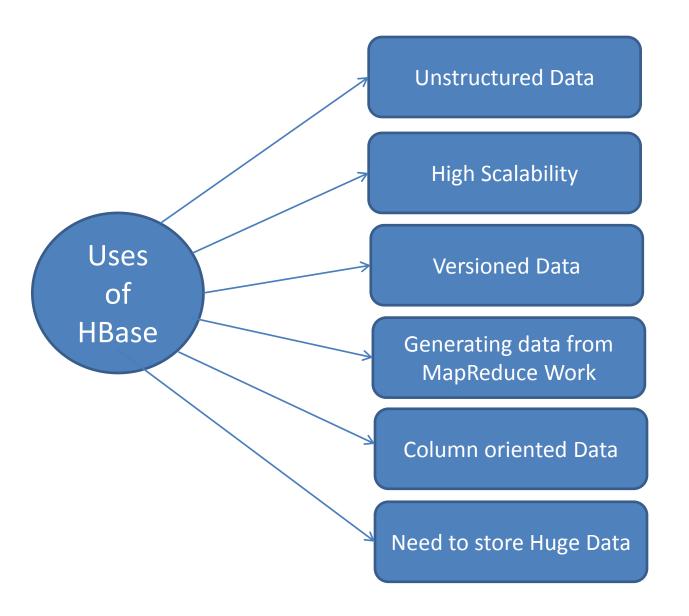


- Consistent Commits are atomic across the entire DFS
- Available Remain Accessible and operated all the time
- Partition Tolerance: Only total network failure can cause system to response incorrectly
  - According to the theorem, a distributed system cannot satisfy all three of these guarantees at the same time

Choose Any Two: CA, CP or AP

# Hbase v/s RDBMS

HBASE	RDBMS
Column Oriented	Row Oriented (Mostly)
Flexible schema, columns can be added in run time	Fixed schema
Good with sparse tables	Not optimized for sparse tables
Joins using MR- not optimizedDesign in such a way as to avoid joinsmillions of columns?	Optimized for Joins
Tight integration with MR	Not really
Scalability - Scale Out Storage (Horizontally)	Scale Up - Difficult to shard and scale
Good for semi-structured data as well as structured data	Good for semi-structured data



#### When **NOT** to use HBase

- ➤ When you have only a few thousands/millions of rows
- ➤ Lacks RDBMS Commands
- When you have hardware less than 5 data nodes when replication factor is 03

### More Hbase Implementation



A number of applications including people search rely on HBase internally for data generation.



Uses HBase to power their Messages infrastructure http://sites.computer.org\debull\A12june\facebook.pdf



We use HBase as a real time data storage and analytics platform.



Uses HBase to store document fingerprint for detecting near-duplications. We have a cluster of few nodes that runs HDFS, mapreduce, and HBase.

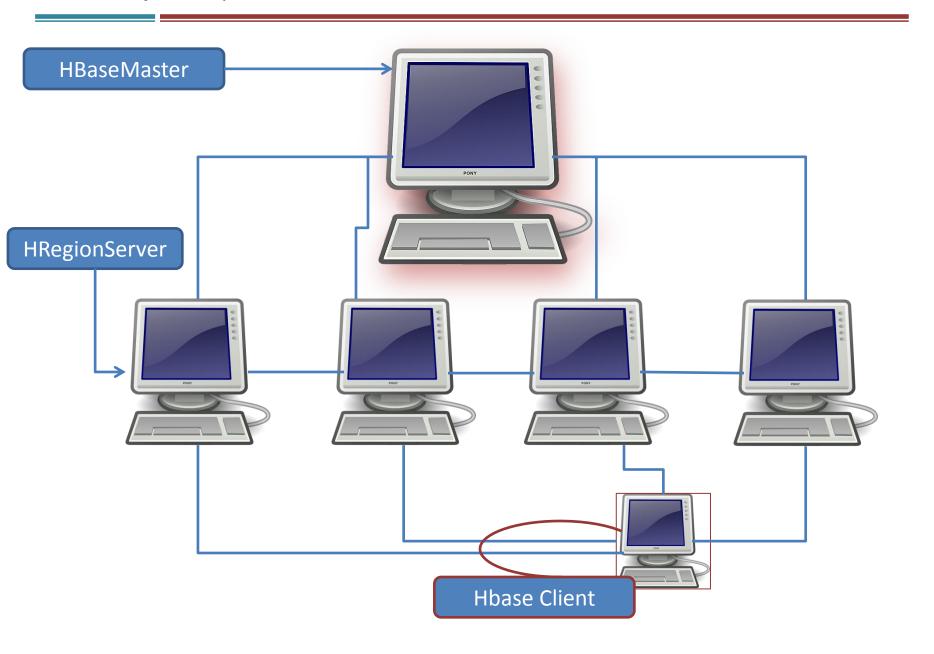


Uses an HBase cluster containing over a billion anonymized clinical records.



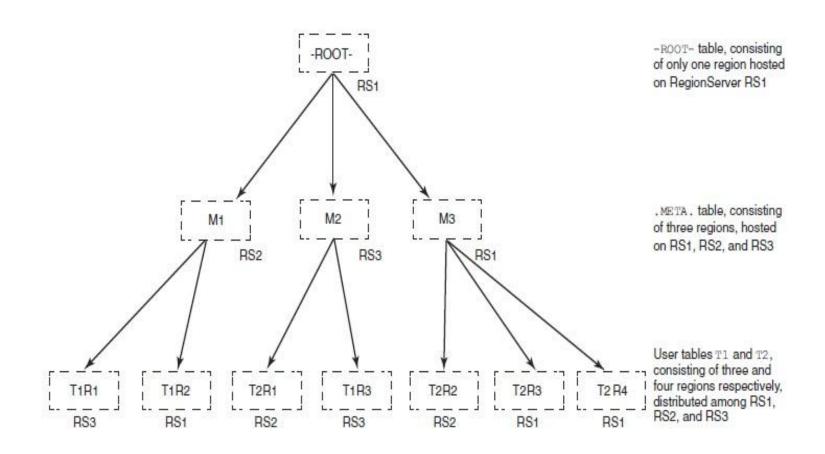
Uses HBase as a foundation for cloud scale storage for a variety of applications.

```
"emp1":{
              "personal":{
                      "name":{
Row Key
                           12312312: "Shyam"
   Col. Family
                      "age": {
   Col. qualifier
                           12312319: "28"
                     "password":{
                           123123232: "password1"
                           123123192: "password2"
String, Hashmap < String, Hashmap < String, String >>>
"emp1", Hashmap<CF-"personal", Hashmap<CQ-"password", Hashmap<string, String>>>
```

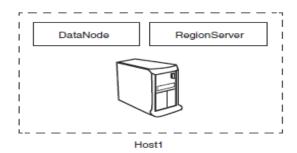


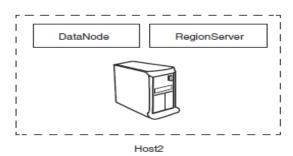
#### When **NOT** to use HBase

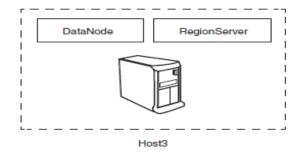
- ➤ When you have only a few thousands/millions of rows
- ➤ Lacks RDBMS Commands
- When you have hardware less than 5 data nodes when replication factor is 03



## How do I find ROOT . META and specific Region – for a particular record







#### T1-R1

00001	John	415-111-1234
00002	Paul	408-432-9922
00003	Ron	415-993-2124
00004	Bob	818-243-9988

RegionServer 1 (RS1), hosting regions R1 of the user table T1 and M2 of . META .

#### .META. - M2

T1:00009-T1:00012	R3	RS2	
-------------------	----	-----	--

#### T1-R2

00005	Carly	206-221-9123
00006	Scott	818-231-2566
00007	Simon	425-112-9877
00008	Lucas	415-992-4432

#### T1-R3

00009	Steve	530-288-9832
00010	Kelly	916-992-1234
00011	Betty	650-241-1192
00012	Anne	206-294-1298

RegionServer 2 (RS2), hosting regions R2, R3 of the user table and M1 of .META.

#### .META. - M1

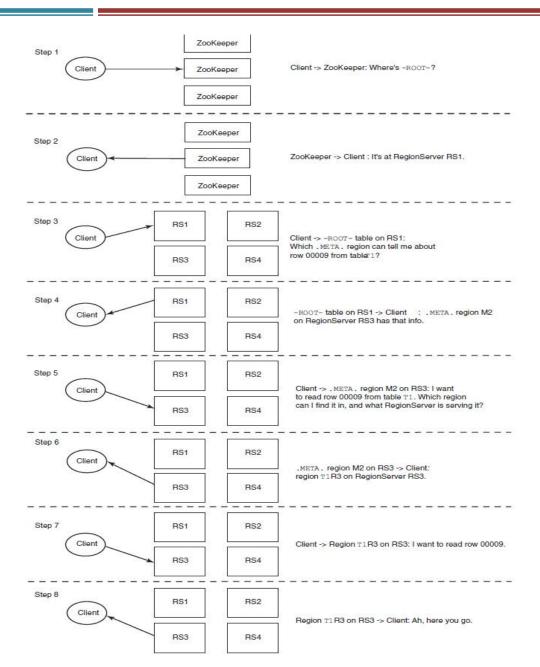
T1:00001-T1:00004	R1	RS1	
T1:00005-T1:00008	R2	RS2	

#### -ROOT-

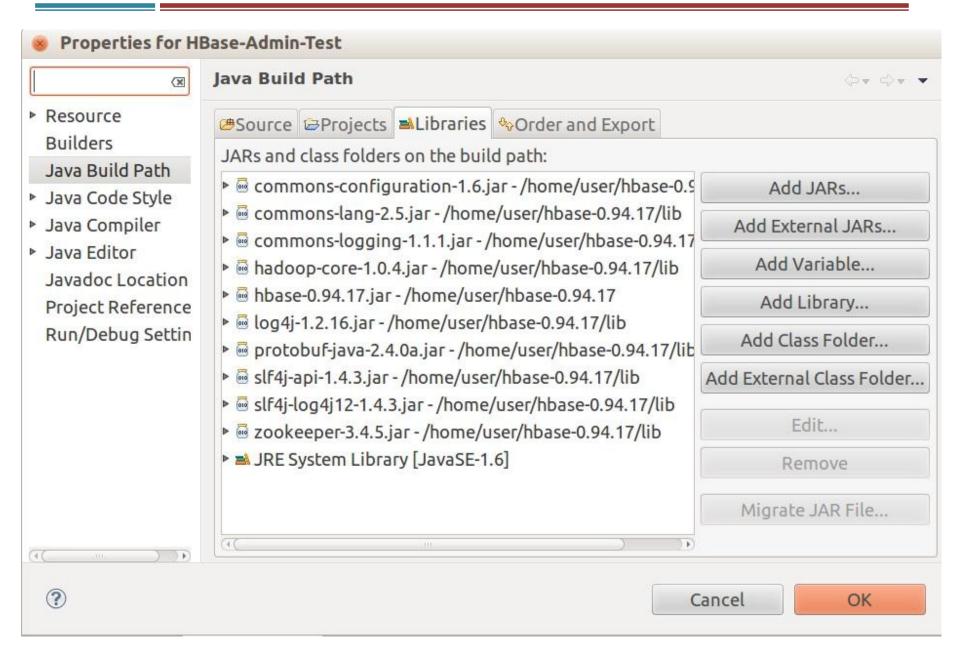
M:T100001-M:T100009	M1	RS2
M:T100010-M:T100012	M2	RS1

RegionServer 3 (RS3), hosting just -ROOT-

# How do I find ROOT. META and specific Region – for a particular record



## List of classes to be included in eclipse before running HBASE program



Bulk Load in Hbase using TSV Format

Before bulk upload, create 'emp' table in hbase with column family as 'cf'

Please ensure to move file 'first.tsv' to move HDFS before following command

Hbase\$ bin/hbase org.apache.hadoop.hbase.mapreduce.lmportTsv

-Dimporttsv.columns=HBASE\_ROW\_KEY, cf:fn, cf:ln emp /first.tsv